

What Does “ $N\times$ Fewer Tokens” Measure? Baselines, Sampling, and Output Form in Repository-Navigation Benchmarks

Anonymous Author(s)
Anonymized for review

Abstract—Repository-graph tools for coding agents are promoted with large token-reduction ratios. We ask how much such a ratio depends on how it is measured. We contribute a fully automatic, model-free benchmark with three label layers that never treat an evaluated tool as ground truth: seeded fixtures whose call sites are known by construction, 17 public Python and TypeScript repositories with independently sampled query targets, and calls observed while running the Python repositories’ test suites. We apply it to two Tree-sitter code-graph tools (one in two versions), ripgrep, and two language servers. A published-style method reproduces its headline on public code: graph answers are $85\times$ smaller than the whole files they replace. The same graph outputs are only $3.0\times$ smaller than `rg -w` output, and $1.1\text{--}2.3\times$ larger than language-server results printed one location per line. Sampling targets from the graph’s own edges raises graph tools’ answer rates from 36–59% to 84–100%, and their recall on observed calls by 0.15–0.24. At matched completeness on held-out fixtures, the most complete graph tool used as many tokens as ripgrep (ratio 1.03, 95% CI 0.94–1.13). We release the harness and data, and recommend that efficiency claims report baseline, sampling frame, output form, and completeness together.

I. INTRODUCTION

Coding agents spend much of their context budget locating code, and graph-based navigation tools promise to answer structural questions such as “who calls X ?” from a pre-built index [1], [2], [3], [4]. Such tools are often promoted with a token ratio. A ratio, however, has three hidden parameters that are not properties of the tool: the *baseline* (whole files, grep, or a language server), the *query sample* (often symbols the graph already connects), and the *output form* (names only, locations, or source). Our starting point was a research prototype, G1, whose own benchmark reported $85.8\times$ fewer tokens than reading the referencing files, measured on graph-selected symbols.

We ask: **Q1** How does the baseline change measured payload ratios? **Q2** How does sampling targets independently of the graph change measured coverage and rankings? **Q3** Under the same output form, and with labels no evaluated tool produced, which tools reach useful token–quality operating points?

Contributions: (1) an open benchmark harness with three independent label layers, automatic validity gates, and output-form control; (2) evidence that, for identical graph outputs, the measured ratio spans more than two orders of magnitude depending on baseline and output form; (3) evidence that

graph-conditioned sampling substantially overstates graph-tool coverage and recall; (4) pattern-level correctness profiles, and recall on runtime-observed calls, for graph, lexical, and language-server retrieval.

II. BENCHMARK DESIGN

Tasks. T1 (primary): given a target declaration, return its direct call sites. The gold unit is a call expression whose callee statically binds to the target. T2: given a call site, return the declaration.

Layer A: seeded fixtures. A generator emits small, executable Python and TypeScript repositories. Each contains one instance of every pattern: direct calls; calls through an import alias, a module or namespace attribute, a re-export, or a typed or super receiver; constructor calls; module-level calls; a same-named method on another class; local shadowing; callbacks; string-dispatched calls; decorators or arrow functions; and comment and string distractors. Half of the fixtures use common, reused names (“ambiguous”). The generator writes a manifest of every expected declaration and site. Python fixtures execute, and TypeScript fixtures compile under `-strict` and run. For every Python target, the call lines observed at runtime equal the manifest’s sites. This gate caught three manifest omissions before any held-out data existed. We use 4+4 calibration and 16+16 held-out fixtures: 224 targets, 160 with explicit call sites.

Layer B: natural repositories. We use 9 Python and 8 TS/JS libraries, selected by fixed criteria before results were observed, plus 3 pilot repositories excluded from all results. Declarations are enumerated by tools that share no code with any graph: Python `ast` and the TypeScript compiler API. *S-ind*, the primary sample, is drawn from a random pool of ≤ 300 non-test declarations, stratified by kind and by the language server’s call fan-out, with inclusion weights recorded; it has 585 targets. *S-cond*, a diagnostic sample of 331 targets, is drawn from declarations with ≥ 1 incoming G1 edge, which is the original benchmark’s rule. Layer B supports agreement, availability, size, and failure rates, but not accuracy.

Layer C: runtime-observed calls. Each Python repository’s full test suite runs under `sys.monitoring`. A call counts only when the callee’s entry is confirmed, since `CALL` fires before invocation. Calls are mapped to the independent index. We report recall on observed calls, a different popu-

TABLE I

Q1: BASELINE/GRAPH PAYLOAD RATIO (REPO-MACRO GEOMETRIC MEAN, S-IND).

Graph tool	whole files	rg-w	LSP JSON	LSP loc.
G1 find_references	131.7	6.32	3.04	0.93
G1 lookup+deps	84.8	3.00	2.18	0.49
codebase-memory current	42.6	3.01	1.93	0.46
codebase-memory v0.5.7	54.4	3.54	1.31	0.44

lation from all calls; unobserved static results are not false positives.

Arms. `rg-w` with 0 and 3 context lines (bare name); `pyright 1.1.414` and the TypeScript 5.9.3 language service (references at the declaration); G1 `find_references`, and G1’s original `lookup_symbol+get_dependencies` pair; and `codebase-memory` [4] at its paper-era release v0.5.7 and at the current main branch. The `codebase-memory` versions are queried through persistent MCP sessions; when a tool reports ambiguity, the follow-up call uses the target’s qualified name and both calls are counted.

Scoring and statistics. Returned facts are mapped to call sites (location tools only), enclosing caller declarations, and files. An answer is *complete* when caller-level recall is 1. Tokens (cl100k) are counted on each tool’s native output, a common one-location-per-line form, and a source-enriched form. An automatic check confirms identical fact sets across forms for every query. Estimates are repository-macro means with 95% hierarchical-bootstrap intervals over repositories, then targets. The primary contrast and a $1.5\times$ margin were frozen before held-out data. One post-freeze adapter fix is reported in §V.

III. RESULTS

A. Q1: the baseline determines the ratio

Fig. 1 and Table I give the ratios. Applying G1’s original definition to the public repositories reproduces its headline: $85.3\times$ as a repository geometric mean ($152.8\times$ pooled, 328 targets), against $85.8\times$ originally reported. That definition uses graph-conditioned targets, whole files selected by the graph, and the `lookup+deps` output. Against `rg-w`, the *same* outputs are $3.0\times$ smaller. Output form alone can flip the sign: the language server’s results are larger than every graph tool’s output as raw JSON, but smaller in the location form (ratios 0.44–0.93). In an exploratory regression, the `rg-w/graph` ratio rises with the number of occurrences of the target’s name in the repository, with an elasticity of ≈ 0.75 (95% CI 0.66–0.88). Lexical retrieval is costly mainly for common, ambiguous names.

B. Q2: graph-conditioned sampling overstates coverage

Table II compares the two samples. Some of the drop in answer rate is legitimate: 39% of S-ind targets have no call sites according to the language server. The recall comparison, however, uses only targets with at least one observed call, so it isolates the effect of conditioning on graph edges. S-cond selects precisely the symbols the graph already connects.

TABLE II

Q2: S-COND $\rightarrow$ S-IND. NON-EMPTY ANSWERS (%), AND CALLER RECALL ON RUNTIME-OBSERVED CALLS (PYTHON).

Arm	non-empty %	observed-call recall
G1 find_references	100 $\rightarrow$ 47	0.85 $\rightarrow$ 0.65
G1 lookup+deps	99 $\rightarrow$ 36	0.81 $\rightarrow$ 0.57
codebase-memory current	88 $\rightarrow$ 56	0.79 $\rightarrow$ 0.61
codebase-memory v0.5.7	84 $\rightarrow$ 59	0.72 $\rightarrow$ 0.57
language server	97 $\rightarrow$ 81	0.84 $\rightarrow$ 0.73
rg-w	100 $\rightarrow$ 100	0.97 $\rightarrow$ 0.94

Independently sampled declarations are often absent from the graph: 16% of S-ind targets are missing from G1. G1’s coverage of non-test declarations is 91–100% in the Python repositories, but 47–68% in five of the eight TS/JS repositories. The causes are unindexed object-literal arrow functions, and same-named nested declarations that collapse into one node (e.g., 303 test-local classes in one repository).

C. Q3: correctness and tokens at matched quality

Table III and Fig. 2 report the held-out fixtures. For the primary contrast we compute the geometric-mean token ratio of each arm to `rg-w` over targets on which both are complete:

- G1 `find_references`: 0.35 [0.33, 0.37], a saving, but it was complete on only 40% of targets, against 80% for `rg-w`;
- G1 `lookup+deps`: 0.65 [0.62, 0.67], not a meaningful saving under the frozen margin;
- `codebase-memory current`: 1.03 [0.94, 1.13], no difference, and the only graph tool as complete as the baselines (89%);
- `codebase-memory v0.5.7`: 0.58 [0.56, 0.62], a saving at 49% completeness;
- `language-server JSON`: 1.75 [1.65, 1.85], a cost.

Small outputs coincided with incomplete answers.

Failures by pattern are as follows. `rg-w` misses only calls through import aliases. `pyright` also misses them; the TypeScript service does not. Of G1’s 240 missed callers (`lookup+deps`): 96 are wrong resolutions to a same-named local function that shadows the target; 32 are module-level calls, which it never attributes to a caller; and 112 are relations it does not record (re-exports, TS constructor calls, and typed-receiver or `super` calls). Ambiguous names lower G1’s recall from 0.73 to 0.53, but leave `codebase-memory current` unaffected (0.95 vs. 0.94). `codebase-memory v0.5.7` misses the typed-receiver and constructor calls that the current version finds. In T2 (definition lookup), every arm found every definition. Only the language server returned exactly one answer (precision 1.00); the other arms returned same-named candidates (precision 0.68–0.76), at 37–98 tokens.

D. Layer C: recall on observed calls

On 323 Python targets with observed calls, caller recall was: `rg-w` 0.95 [0.92, 0.98]; language server 0.79 [0.64, 0.90]; G1 0.75 and 0.69; `codebase-memory` 0.70 (current) and 0.64




Fig. 1. Q1. Payload of each baseline relative to each graph tool’s native output, for the same queries (S-ind, non-empty graph answers). Points are repository geometric means; bars are the mean across repositories. The ratio moves by over two orders of magnitude with the choice of baseline and output form.




Fig. 2. Q3. Held-out fixtures, Python and TypeScript pooled: native output tokens vs. caller-level recall per arm (repository means; G1 pair = lookup+deps). Small outputs coincide with low recall; the only graph tool with baseline-level recall needs baseline-level tokens.

TABLE III
Q3: HELD-OUT FIXTURES, T1, CALLER LEVEL, 160 TARGETS.

Arm	recall	precision	complete %	tokens
rg -w	0.95	0.35	80.0	140
rg -w -C3	0.95	0.35	80.0	465
language server	0.97	0.64	90.0	236
G1 find_references	0.63	0.68	40.0	50
G1 lookup+deps	0.47	1.00	40.0	94
codebase-memory current	0.94	0.98	89.4	160
codebase-memory v0.5.7	0.71	0.90	49.4	70

(v0.5.7). The language server’s misses concentrate in repositories that ship `.pyi` stubs, because pyright binds public-API calls to the stub declarations rather than the implementation.

IV. RECOMMENDATIONS

- Report the baseline, and more than one: the same graph outputs measured $0.44\times$ to $134\times$.
- Sample targets independently of the tool under test: edge-conditioned sampling raised graph tools’ answer rates by

25–64 points.

- Count tokens in a common output form, and state what each form contains: names-only outputs carry less information.
- Report completeness next to tokens: at matched completeness, the most complete graph tool saved nothing over rg -w.
- Do not equate payload tokens with agent cost. Agent behaviour and prompt caching are out of scope here, and agent-level studies [4], [5] measure something different.

V. THREATS TO VALIDITY

Construct. Payload tokens are not agent cost [9], [10]. Static possible-target labels encode one semantics. Caller-level scoring credits any location inside a true caller.

Internal. After freezing, we found that pyright answers before its workspace is indexed. We added a readiness wait (0/80 reference sets changed on re-query 30s later) and re-ran all Python jobs. This raised layer-C language-server recall from 0.73 to 0.79 and changed no conclusion; prefix results are released. Before freezing, name-resolution rules were corrected so that tool outputs are not penalised by our own matching. Unresolved output (0–6%) still counts against the tool. Tokens were counted with gpt-tokenizer’s bundled cl100k ranks, validated against known token ids.

External. Fixtures are cleaner than real code. The repositories are a purposive selection, and only two languages are covered. Layer C covers only exercised code. There is no human audit, so accuracy on natural code is not measured.

VI. RELATED WORK

Code graphs for agents improve repository-level tasks [1], [2], [3]. Codebase-Memory [4] reports about $10\times$ fewer agent tokens than an agent using file reading and grep, at somewhat lower answer quality (83% vs. 92%). Xu [5] finds that language servers seldom save agent tokens. CodeNib [6] compares static and live navigation on graph-derived queries.

Lexical search is a strong agent baseline [7]. PyCG bounds what static call graphs achieve [8]. We complement these studies with a model-free, cross-tool measurement whose labels are independent of every evaluated tool, and we quantify how sampling and output form change conclusions.

VII. DATA AVAILABILITY

Harness, fixtures generator, pinned corpus commits, raw per-query outputs, traces, analysis, and the frozen protocol: <https://anonymous.4open.science/r/XXXX> (anonymized).

REFERENCES

- [1] S. Ouyang *et al.*, “RepoGraph: Enhancing AI software engineering with repository-level code graph,” in *Proc. ICLR*, 2025.
- [2] X. Liu *et al.*, “CodexGraph: Bridging large language models and code repositories via code graph databases,” in *Proc. NAACL*, 2025.
- [3] Z. Chen *et al.*, “LocAgent: Graph-guided LLM agents for code localization,” in *Proc. ACL*, 2025.
- [4] M. Vogel *et al.*, “Codebase-Memory: Tree-sitter-based knowledge graphs for LLM code exploration via MCP,” arXiv:2603.27277, 2026.
- [5] P. Xu, “Does a language server save tokens for coding agents? A measurement methodology and preliminary study,” arXiv:2608.13568, 2026.
- [6] “CodeNib: A multi-view data system for serving repository context to coding agents,” arXiv:2607.25431, 2026.
- [7] S. Sen *et al.*, “Is grep all you need? How agent harnesses reshape agentic search,” arXiv:2605.15184, 2026.
- [8] V. Salis *et al.*, “PyCG: Practical call graph generation in Python,” in *Proc. ICSE*, 2021.
- [9] L. Bai *et al.*, “How do AI agents spend your money? Analyzing and predicting token consumption in agentic coding tasks,” arXiv:2604.22750, 2026.
- [10] S. Weinberger and A. Hozes, “Token reduction is not cost reduction: An empirical study of end-to-end efficiency in API-based coding agents,” arXiv:2607.12161, 2026.